



FACULTY OF SCIENCES OF THE UNIVERSITY  
OF LISBON

# Advanced Machine Learning Assignment 1 - Backpropagation

*55761 · António Rebelo*  
*55877 · Marco Mercier*

taught by  
Professor Luis CORREIA

February 15, 2021

## 1 Problem 1

a) The equation for the local field of the first neuron ( $v_1$ ) is:

$$v_1 = w_{11}x_1 + w_{12}x_2 + b_1 = x_1 + x_2 - 1.5$$

For the second neuron ( $v_2$ ):

$$v_2 = w_{21}x_1 + w_{22}x_2 + w_{23}\phi(v_1) + b_2 = x_1 + x_2 - 2\phi(v_1) - 0.5$$

Thus, the output will be:

$$O = \phi(x_1 + x_2 - 2\phi(x_1 + x_2 - 1.5) - 0.5)$$

Where:

$$\phi(n) = \begin{cases} 0 & n \leq 0 \\ 1 & n > 0 \end{cases}$$

is the activation function. The separation planes will be where  $n = 0$ , which is when the activation function makes the neuron fire differently.

$$v_1 = 0 \Leftrightarrow x_1 = 1.5 - x_2$$

$$v_2 = 0 \Leftrightarrow x_1 = 0.5 + 2\phi(v_1) - x_2$$

Since  $x_1, x_2 \in \{0, 1\}$ , lets explore for both values:

$$\begin{aligned} x_1 = 0 &\Rightarrow \begin{cases} N_1 & x_2 = 1.5 \\ N_2 & x_2 = 0.5 \end{cases} \\ x_2 = 0 &\Rightarrow \begin{cases} N_1 & x_1 = 1.5 \\ N_2 & x_1 = 0.5 \end{cases} \end{aligned}$$

Thus, the separation planes is:

$$\begin{cases} x_2 = -x_1 + 1.5 \\ x_2 = -x_1 + 0.5 \end{cases}$$

b) The Truth table is obtained by using the output equation above:

$x_0$	$x_1$	Output
0	0	0
0	1	1
1	0	1
1	1	0

c) Considering  $a = 0$  for the activation function  $\phi(v) = \sigma(v) = \frac{1}{1+e^{-v}}$ , learning rate  $\eta = 1$ , the initial weights to be the ones used before and that the previous Truth Table provides the desired outputs.

The delta rule is  $\Delta w_{ji} = \eta \delta_j y_i$ , where  $\delta_j$  is the local gradient of neuron  $j$  and  $y_i$  the output of the neuron  $i$ . Using the sigmoid activation function, the local gradient of an output neuron  $k$  is:

$$\delta_k = y_k(1 - y_k)(d_k - y_k)$$

and the local gradient for a hidden layer neuron  $j$  is:

$$\delta_j = y_j(1 - y_j) \sum_k \delta_k w_{kj}$$

Lets run the first iteration of the first epoch for  $(x_1, x_2, O) = (0, 0, 0)$ :

$$v_1 = -1.5 \Rightarrow y_1 = \phi(v_1) = \sigma(v_1) = 0.182$$

$$v_2 = -0.5 - 2y_1 = -0.864 \Rightarrow y_2 = 0.297$$

$$\delta_2 = 0.297(1 - 0.297)(0 - 0.297) = -0.062$$

Lets adjust the weights of the output neuron:

$$w_{21} = w_{21} + \eta \delta_2 x_1 = 1 + 1 * -0.062 * 0 = 1$$

$$w_{22} = w_{22} + \eta \delta_2 x_2 = 1 + 0 = 1$$

$$w_{23} = w_{23} + \eta \delta_2 y_1 = -2 + 1 * -0.062 * 0.182 = -2.011$$

$$b_2 = b_2 + \eta \delta_2 = -0.5 + 1 * -0.062 = -0.562$$

Now for the hidden layer neuron:

$$\delta_1 = y_1(1 - y_1) \delta_2 w_{23} = 0.182(1 - 0.182) * -0.062 * -2 = 0.0185$$

$$w_{11} = w_{11} + \eta \delta_1 x_1 = 1 + 1 * 0.0185 * 0 = 1$$

$$w_{12} = w_{12} + \eta \delta_1 x_2 = 1 + 1 * 0.0185 * 0 = 1$$

$$b_1 = b_1 + \eta \delta_1 = -1.5 + 1 * 0.0185 = -1.482$$

# Hand-in-1\_ex3

February 14, 2021

```
[1]: import torch.nn as nn
import torch
import torch.nn.functional as F
from torch.autograd import Variable
import numpy as np
import random
import pandas as pd
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
%matplotlib inline
```

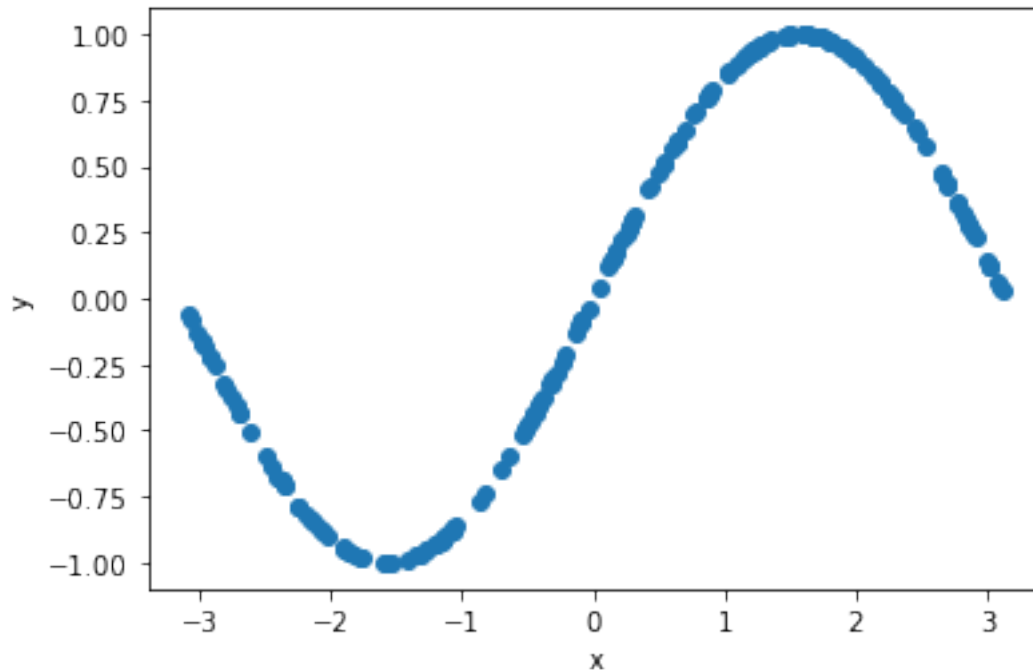
We start by defining x and y as is requested

```
[2]: #to guarantee reproducible results
np.random.seed(0)

x = [random.uniform(-np.pi,np.pi) for _ in range(200)]
y = np.sin(x)
```

```
[3]: #Visualizing the data:
fig, ax = plt.subplots()
ax.scatter(x,y);
ax.set_xlabel('x')
ax.set_ylabel('y')
```

```
[3]: Text(0, 0.5, 'y')
```



```
[4]: #pytorch only trains on its special Variable class so we need to convert the
      ↪variables
x = Variable(torch.tensor(x))
y = Variable(np.sin(x))
```

```
[5]: #the dimensions of the vectors also need to be adjusted because of how pytorch
      ↪multiplies matrices to train NN's
x = x.unsqueeze(dim=-1)
y = y.unsqueeze(dim=-1)
```

We split the data into training, test and validation set

```
[6]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.
      ↪15, random_state = 20)
```

```
[7]: x_train, x_val, y_train, y_val = train_test_split(x_train, y_train, test_size=0.
      ↪17, random_state = 20)
```

## 0.1 Setting up the network

```
[8]: #set the random seed to initialize weights
torch.manual_seed(1);
```

```
[9]: #Here comes the network with one hidden layer of two neurons (nn.Linear(1,2))
      → and a tanh activation function
class NN(torch.nn.Module):
    def __init__(self):
        super(NN,self).__init__()
        self.layer1 = nn.Linear(1,15)
        self.predict = nn.Linear(15,1)

    def forward(self,x):
        x = torch.tanh(self.layer1(x))
        x = self.predict(x)
        return x
```

## 0.2 Training the model

```
[10]: #instantiate the network
Net = NN()
#set up the optimizer as Adam with learning rate 0.2
opt = torch.optim.Adam(Net.parameters(), lr = 0.2)
#set up the los function as MSE
loss_func = torch.nn.MSELoss()
```

```
[11]: #now the actual training
      #we set the number of training epochs
n_epochs = 300

for epoch in range(n_epochs):

    #calculate predictions
    preds = Net(x_train)

    #calculate loss
    loss = loss_func(preds,y_train)

    #set gradients to zero
    opt.zero_grad()
    #backpropagate from the loss
    loss.backward()
    #update weights
    opt.step()
```

### 0.3 Results on validation set

Note: in a real world problem, this would be overfitting, but in a academic setting such as this one, that has no meaning.

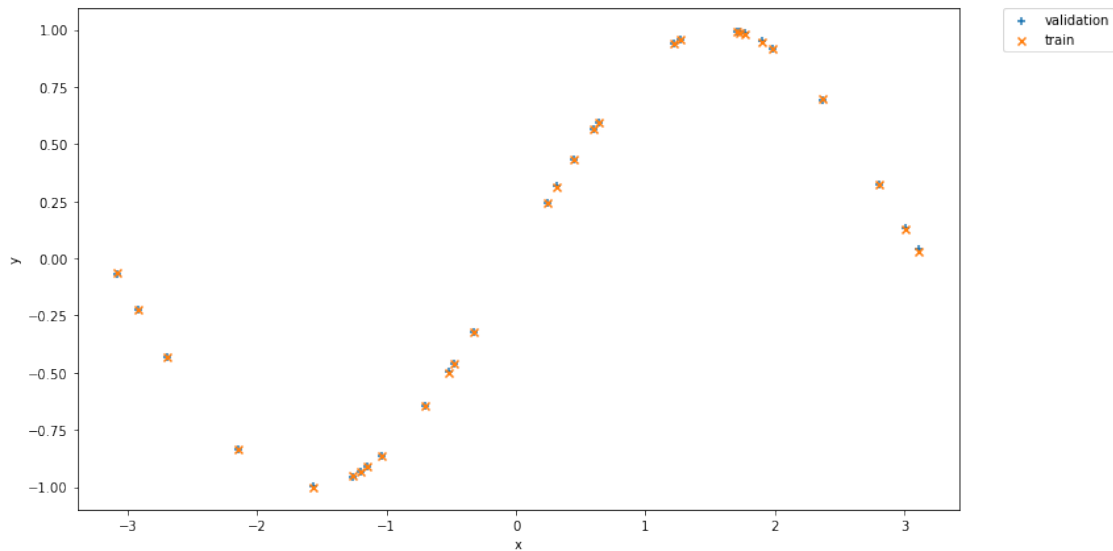
```
[12]: #the predictions for the validation set:
preds_val = Net(x_val)

#the MSE
loss_func(preds_val,y_val)
```

```
[12]: tensor(1.0091e-05, grad_fn=<MseLossBackward>)
```

The loss is very very low so the model is fitting very well. Checking the results graphically:

```
[13]: fig, ax = plt.subplots(figsize=(12,7))
ax.scatter(x_val.flatten().detach().numpy(),preds_val.flatten().detach().
    ↪numpy(), label = "validation",marker="+");
ax.scatter(x_val.flatten().detach().numpy(),y_val.flatten().detach().numpy() ,
    ↪label = "train",marker="x");
ax.set_xlabel('x')
ax.set_ylabel('y')
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left', borderaxespad=0.);
```



```
[ ]:
```

## 0.4 Results on test set

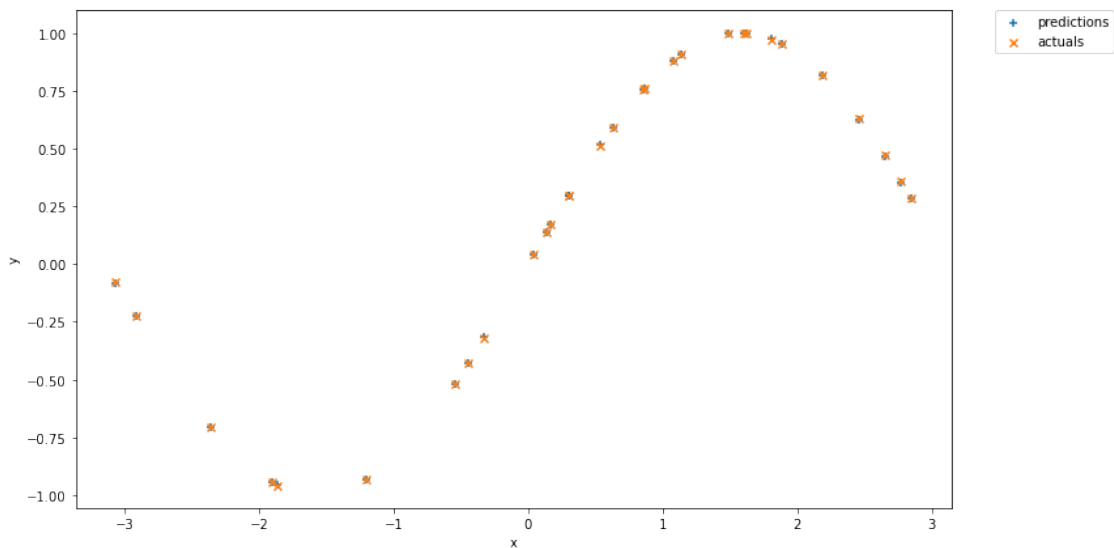
```
[14]: #the predictions for the test set:
      preds = Net(x_test)

      #the MSE
      loss_func(preds,y_test)
```

```
[14]: tensor(5.2821e-06, grad_fn=<MseLossBackward>)
```

The loss is very very low so the model is fitting very well. Checking the results graphically:

```
[15]: fig, ax = plt.subplots(figsize=(12,7))
ax.scatter(x_test.flatten().detach().numpy(),preds.flatten().detach().numpy(),
↪ label = "predictions",marker="+");
ax.scatter(x_test.flatten().detach().numpy(),y_test.flatten().detach().numpy()
↪ , label = "actuals",marker="x");
ax.set_xlabel('x')
ax.set_ylabel('y')
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left', borderaxespad=0.);
```



[ ]:

[ ]:

[ ]: